

# Research paper presentation: “De-indirection for Flash-based SSDs with Nameless Writes”

Y. Zhang, L. P. Arulraj, A. C. Arpaci-Dusseau, R. H. Arpaci-Dusseau

Federico Wasserman & Rodolphe Lepigre

MOSIG - Parallel, Distributed and Embedded Systems

December 18, 2012

# Outline

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

## What are Nameless Writes?

- New device interface for SSDs
- Remove the need for indirection
- Idea: the device chooses WHERE to write

## How are Nameless Writes different?

Usual Writes:

- The FS requests the writing of data at some location
- The device performs the write

Nameless Writes:

- The FS requests the writing of data
- The device performs the write
- Address returned to the FS

- 1 Introduction
- 2 SSD principles**
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion



- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs**
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion



# SSDs need indirection

- Indirection is used to implement wear-leveling
- Absolutely necessary to ensure reasonable lifetime
- Problem: need to store indirection table
- 3 main techniques:
  - Full-page mapping
  - Block mapping
  - Hybrid mapping

## Full-page mapping

- Each page can be mapped
- Consider 32-bit pointers per 2KB pages
- With 1TB SSD, 2GB indirection table
- Problem: Great space overhead, DRAM is expensive

## Block mapping

- Mapping at block-level (128 pages)
- 32MB indirection table in the same settings
- Smaller memory overhead
- Problem: high garbage collection cost (Gupta et al.)

## Hybrid mapping

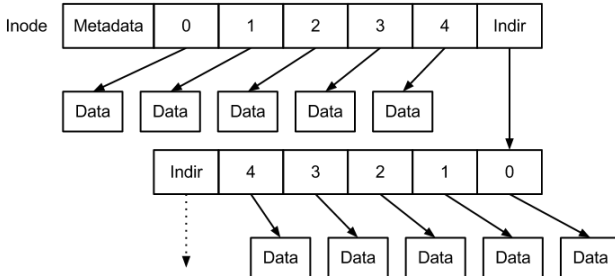
- Map most data at block level
- Small page-mapped area
- Keeps space overhead low
- Avoids garbage collection overhead
- Problem: garbage collection can still hurt performances
- Problem: very complex FTL (Flash Translation Layer)
- Solution: Nameless Writes

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes**
- 5 Evaluation
- 6 Conclusion


## Main interface

```
Nameless_Write(data , len) : phys@  
Nameless_Overwrite(phys@ , data , len) : new@  
Physical_Read(phys@ , len) : data  
Free(vitr/phys@ , len)
```

## Inode, and file structure



## File tree





## Virtual read / write interface

```
Virtual_Write(virt@ , data , len)  
Virtual_Read(virt@ , len) : data
```

## Migration callback

Migration [ Callback ] ( old\_phys@ , new\_phys@ )

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation**
- 6 Conclusion



- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion**

