

Research paper presentation: “De-indirection for Flash-based SSDs with Nameless Writes”

Y. Zhang, L. P. Arulraj, A. C. Arpaci-Dusseau, R. H. Arpaci-Dusseau

Federico Wasserman & Rodolphe Lepigre

MOSIG - Parallel, Distributed and Embedded Systems

December 19, 2012

Outline

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

What are Nameless Writes?

- New device interface for SSDs
- Remove the need for indirection
- Idea: the device chooses WHERE to write

How are Nameless Writes different?

Usual Writes:

- The FS requests the writing of data at some location
- The device performs the write

Nameless Writes:

- The FS requests the writing of data
- The device performs the write
- Address returned to the FS

- 1 Introduction
- 2 SSD principles**
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

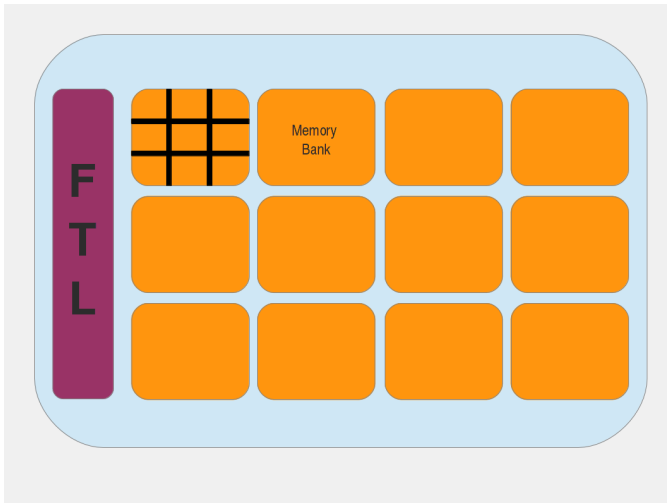
SSDs are not Hard Drives

- Essentially different from Hard Drives
- Fast constant access to a random position
- Limited number of writes


SSDs Internally (1) - General View



SSDs Internally (2) - Blocks




SSDs Internally (3) - Blocks and Pages




Writing to a block

- Page needs to be in a erased state
- Cannot overwrite data
- Erasing is done at block level



Overwrite - Erase/Program



Wear Leveling and Garbage Collection

- Reusing same block reduces its lifetime
- Erase/Program is expensive
- Move data to better use blocks
- Garbage Collection - Recover invalid pages

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs**
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion

SSDs need indirection

- Indirection is used to implement wear-leveling
- Absolutely necessary to ensure reasonable lifetime
- Problem: need to store indirection table
- 3 main techniques:
 - Full-page mapping
 - Block mapping
 - Hybrid mapping

Full-page mapping

- Each page can be mapped
- Consider 32-bit pointers per 2KB pages
- With 1TB SSD, 2GB indirection table
- Problem: Great space overhead, DRAM is expensive

Block mapping

- Mapping at block-level (128 pages)
- 32MB indirection table in the same settings
- Smaller memory overhead
- Problem: high garbage collection cost (Gupta et al.)

Hybrid mapping

- Map most data at block level
- Small page-mapped area
- Keeps space overhead low
- Avoids garbage collection overhead
- Problem: garbage collection can still hurt performances
- Problem: very complex FTL (Flash Translation Layer)
- Solution: Nameless Writes

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes**
- 5 Evaluation
- 6 Conclusion

Reminder

Idea: the device chooses WHERE to write

- The FS requests the writing of data
- The device performs the write
- Address returned to the FS

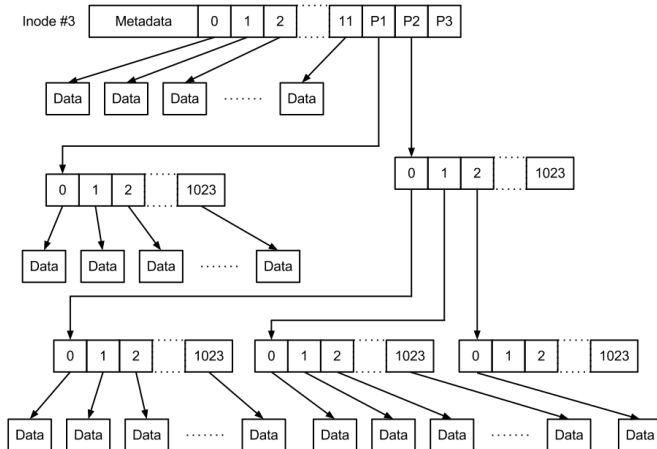
Main interface

```
Nameless_Write(data , len) : phys@  
Nameless_Overwrite(phys@ , data , len) : new@  
Physical_Read(phys@ , len) : data  
Free(vitr/phys@ , len)
```


Recursive update problem

- Problem with this interface: recursive update
- File modification imply inode update
- The inode will move (Nameless overwrite)
- Every structure pointing to it will have to be updated
- ...

Inode, and file structure



File tree



Solution: Segmented address space

- Large physical address space (for nameless writes)
- Small virtual address space (for traditional writes)
- Idea: keep pointer-based structures in virtual space

Segmented address space

Virtual read / write interface

```
Virtual_Write(virt@ , data , len)  
Virtual_Read(virt@ , len) : data
```

Migration callback

- Callback provided for the SSD to notice data migration to the FS
- Useful for it to reclaim blocks (garbage collection)

Migration [Callback] (old_phys@ , new_phys@)

- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation**
- 6 Conclusion


Setup

- Emulator with three different FTLs
 - Page-Level,
 - Hybrid 5% page-level mapping
 - Nameless Writes

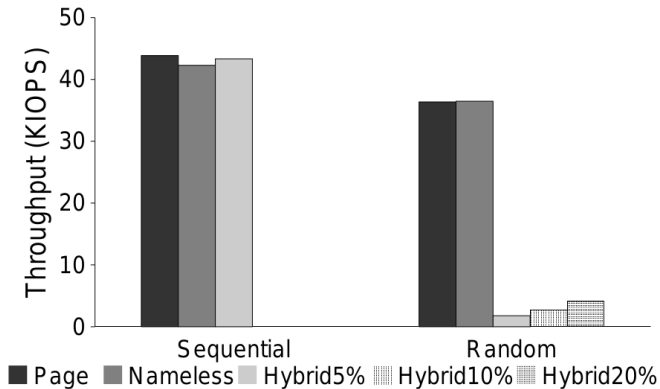
Memory Consumption

Image Size	Page	Hybrid	Nameless
328 MB	328 KB	38 KB	2.7 KB
2 GB	2 MB	235 KB	12 KB
10 GB	10 MB	1.1 MB	31 KB
100 GB	100 MB	11 MB	251 KB
400 GB	400 MB	46 MB	1 MB
1 TB	1 GB	118 MB	2.2 MB



Performance



Performance



Performance



- 1 Introduction
- 2 SSD principles
- 3 Indirection in SSDs
- 4 Nameless Writes
- 5 Evaluation
- 6 Conclusion**

Conclusion

- New write interface
- Improves random-write performance
- Reduce space costs
- Future work: Port to other file systems

Questions

- Questions?